

PARESH PUROHIT

+91 6377069019 · pareshpurohit1920@gmail.com · Vadodara, Gujarat · [GitHub](#) · [LinkedIn](#)

PROFESSIONAL SUMMARY

Final-year B.Tech Computer Science student (May 2026) specialising in Data Science and AI, with hands-on experience building ML pipelines, computer vision systems, IoT-applicable predictive models, and LLM-powered applications. Demonstrated results include 97%+ model accuracy, 40% reduction in search latency, and 30% faster model training. Proficient in Python, Scikit-learn, XGBoost, TensorFlow, OpenCV, LangChain, Flask, and SQL. Experienced in end-to-end data pipeline and ETL development, EDA, feature engineering, and REST API deployment. Passionate about applying AI and data science to IoT, home automation, and real-world business workflows — including demand forecasting, predictive maintenance, and customer analytics. Actively exploring NLP, BI dashboards (Tableau/Looker), and Edge AI (TensorFlow Lite, ONNX) to deliver production-grade solutions.

TECHNICAL SKILLS

Languages: Python, SQL, JavaScript, C++

ML & Data: Scikit-learn, XGBoost, NumPy, Pandas, SciPy, Feature Engineering, EDA, Statistical Modeling, Demand Forecasting, Predictive Maintenance, Lead Scoring, Pipeline Automation, RAG Pipeline

Deep Learning / CV: TensorFlow, Keras, PyTorch (basic), CNN, RNN, Transfer Learning, OpenCV, YOLOv8, Object Detection, Edge AI (TensorFlow Lite, ONNX)

NLP: NLTK, SpaCy, Hugging Face, Text Preprocessing, Sentiment Analysis, LangChain, LLMs, OpenAI API

Databases: MySQL, MongoDB, FAISS Vector Store

APIs & Deployment: Flask, FastAPI, REST API, Docker, AWS (S3, EC2), MLOps, Streamlit, Model Deployment, Postman

BI & Visualization: Matplotlib, Seaborn, Tableau (basic), Google Looker Studio (basic)

IoT & Protocols: MQTT (basic), JSON log parsing, IoT data pipelines, Edge AI deployment

Tools: Git, GitHub, Jupyter, VS Code

EDUCATION

Bachelor of Technology (B.Tech) — Computer Science Engineering Aug 2022 – May 2026

Parul University, Vadodara · Specialisation: Data Science & AI

Relevant Coursework: Data Structures & Algorithms, DBMS, Machine Learning, Computer Vision, Statistics & Probability, Linear Algebra

INTERNSHIP EXPERIENCE

Data Science Intern | CODTECH IT Solutions

Jan 2026 – Apr 2026

- Built and evaluated predictive models (Linear Regression, Decision Trees, XGBoost) on 50,000+ record datasets, achieving average R^2 of 0.87 — directly applicable to demand forecasting and lead scoring use cases.
- Designed and automated ETL/data pipeline workflows in Python and Pandas, reducing data inconsistencies by ~35% and improving downstream model reliability.
- Created reporting dashboards to monitor data quality KPIs, enabling faster QA decisions and cutting manual validation time by ~20%.
- Maintained reproducible ML pipelines in Jupyter Notebooks with Git version control, enabling smooth collaboration across a 4-member team.

PROJECTS

Fake Currency Detection System

[Github](#)

[LiveDemo](#)

| Python · OpenCV · Scikit-learn · TensorFlow · Keras · CNN · YOLOv8

- Built a computer vision pipeline using Python, OpenCV, and TensorFlow that identifies counterfeit currency with 97%+ accuracy across 2,000+ test images — deployable in real-world and IoT edge environments (TF Lite compatible).
- Reduced model training time by 30% by automating image preprocessing (resizing, noise reduction) and parallelising data augmentation steps.

- Benchmarked CNN, SVM, and Random Forest models; selected and deployed the CNN variant for its superior precision-recall balance on imbalanced class data.

DocMind — AI Document Assistant

[Github](#) [LiveDemo](#)

| *Python · Streamlit · LangChain · FAISS · OpenAI API · LLMs · RAG Pipeline*

- Designed and deployed an interactive RAG-based web application on Hugging Face Spaces enabling users to query uploaded PDF documents using Large Language Models and OpenAI API integration.
- Engineered an end-to-end data pipeline for document intelligence: text chunking → embeddings → FAISS vector store (LangChain), reducing document search latency by 40% vs. keyword-search baseline.
- Integrated NLP text preprocessing, sentiment-aware chunking, and REST-style callbacks via Streamlit front-end, demonstrating full-stack AI application deployment applicable to IoT and business workflows.

CERTIFICATIONS & CONTINUOUS LEARNING

- Python for Everybody — Google on Coursera [\(Verified Certificate\)](#)
- Actively follows AI/ML research via arXiv, Hugging Face Blog, and Papers with Code; participates in Kaggle competitions to stay current with industry techniques including NLP, computer vision, and LLM fine-tuning.